

Our focus in this lecture is more on Kubernetes-related security. What are the risks and what measures do you need to take to secure the cluster?

Kube API server

is at the center of all operations within Kubernetes.

We interact with it through the kube control utility or by accessing the API directly.

And through that, you can perform almost any operation on the cluster.

So that's the first line of defense, controlling access to the API server itself.

We need to make two types of decisions.

Who can access the cluster, and what can they do?

Who can access the API server

is defined by the authentication mechanisms.

There are different ways that you can authenticate to the API server, starting with user IDs

and passwords stored in static files or tokens,

certificates, or even integration

with external authentication providers like LDAP.

Finally, for machines, we create service accounts.

We will look at these in more detail in the upcoming lectures.

Once they gain access to the cluster, what can they do is defined by authorization mechanisms.

Authorization is implemented using role-based access controls where users are associated to groups with specific permissions.

In addition, there are other authorization modules like the attribute-based access control, node authorizers, web pods, et cetera.

Again, we look at these in more detail in the upcoming lectures.

All communication with the cluster

between the various components such as the etcd cluster, the kube controller, manager scheduler, API server, as well as those running on the worker nodes such as the kubelet and the kube-proxy is secured using TLS encryption.

We have a section entirely for this where we discuss and practice how to set up certificates between the various components.

What about communication

between applications within the cluster?

By default, all ports can access all other ports within the cluster.

Now, you can restrict access between them using network policies.

Instructor: Hello, and welcome to this lecture on authentication in a Kubernetes cluster.

As we have seen already, the Kubernetes cluster consists of multiple nodes, physical or virtual, and various components that work together.

We have users like administrators that access the cluster to perform administrative tasks, the developers that access the cluster to test or deploy applications, we have end users who access the applications deployed on the cluster, and we have third party applications accessing the cluster for integration purposes.

Throughout this section, we will discuss how to secure our cluster by securing the communication between internal components and securing management access to the cluster through authentication and authorization mechanisms.

In this lecture, our focus is on securing access to the Kubernetes cluster with authentication mechanisms.

So we talked about the different users that may be accessing the cluster.

Security of end users who access the applications deployed on the cluster is managed by the applications themselves internally.

So we will take them out of our discussion.

Our focus is on users access to the Kubernetes cluster for administrative purposes.

So we are left with two types of users.

Humans such as the administrators and developers and robots such as other processes or services or applications that require access to the cluster.

Kubernetes does not manage user accounts natively.

It relies on an external source like a file with user details or certificates or a third party identity service like LDAP to manage these users.

And so you cannot create users in a Kubernetes cluster or view the list of users like this.

However, in case of service accounts, Kubernetes can manage them.

You can create and manage service accounts using the Kubernetes API.

We have a section on service accounts exclusively where we discuss and practice more about service accounts. For this lecture, we will focus on users in Kubernetes.

All user access is managed by the API server whether you're accessing the cluster through kube control tool or the API directly, all of these requests go through the kube-apiserver. The kube-apiserver authenticates the request before processing it.

So how does the kube-apiserver authenticate? There are different authentication mechanisms that can be configured.

You can have a list of username and passwords in a static password file, or usernames and tokens in a static token file, or you can authenticate using certificates. And another option is to connect to third party authentication protocols like LDAP, Kerberos, et cetera.

We will look at some of these next.

Let's start with static password and token files as it is the easiest to understand.

Let's start with the simplest form of authentication.

You can create a list of users and their passwords in a CSV file and use that as the source for user information.

The file has three columns: password, username, and user id.

We then pass the file name as an option to the kube-apiserver.

Remember the kube-apiserver service and the various options we looked at earlier in this course, that is where you must specify this option.

You must then restart the kube-apiserver for these options to take effect.

If you set up your cluster using the kubeadm tool, then you must modify the kube-apiserver pod definition file.

The kubeadm tool will automatically restart the kube-apiserver once you update this file.

To authenticate using the basic credentials while accessing the API server, specify the user and password in a curl command like this.

In the CSV file with the user details that we saw, we can optionally have a fourth column with the group details to assign users to specific groups.

Similarly, instead of a static password file, you can have a static token file.

Here, instead of password, you specify a token.

Pass the token file as an option token-auth-file to the kube-apiserver.

While authenticating, specify the token as an authorization header token to your request like this.

That's it for this lecture.

Remember that this authentication mechanism that stores usernames, passwords, and tokens in clear text in a static file is not a recommended approach as it is insecure.

But I thought this was the easiest way to understand the basics of authentication in Kubernetes.

Going forward, we will look at other authentication mechanisms.

I also wanna point out that if you are trying this out in a kubeadm setup, you must also consider volume mounts to passing the auth file.

Details about these are available

in the article that follows,

and remember to set up authorization for the new users.

We will discuss about authorization later in this course.

## Article on Setting up Basic Authentication

Setup basic authentication on Kubernetes (Deprecated in 1.19)

Note: This is not recommended in a production environment. This is only for learning purposes. Also note that this approach is deprecated in Kubernetes version 1.19 and is no longer available in later releases

Follow the below instructions to configure basic authentication in a kubeadm setup.

Create a file with user details locally at /tmp/users/user-details.csv

1. # User File Contents
2. password123,user1,u0001
3. password123,user2,u0002
4. password123,user3,u0003
5. password123,user4,u0004
6. password123,user5,u0005

Edit the kube-apiserver static pod configured by kubeadm to pass in the user details.

The file is located at /etc/kubernetes/manifests/kube-apiserver.yaml

1. apiVersion: v1
2. kind: Pod
3. metadata:
4. name: kube-apiserver
5. namespace: kube-system
6. spec:
7. containers:
8. - command:
9. - kube-apiserver
10. <content-hidden>
11. image: k8s.gcr.io/kube-apiserver-amd64:v1.11.3
12. name: kube-apiserver
13. volumeMounts:
14. - mountPath: /tmp/users

15. name: usr-details
16. readOnly: true
17. volumes:
18. - hostPath:
19. path: /tmp/users
20. type: DirectoryOrCreate
21. name: usr-details

Modify the kube-apiserver startup options to include the basic-auth file

1. apiVersion: v1
2. kind: Pod
3. metadata:
4. creationTimestamp: null
5. name: kube-apiserver
6. namespace: kube-system
7. spec:
8. containers:
9. - command:
10. - kube-apiserver
11. - --authorization-mode=Node,RBAC
12. <content-hidden>
13. - --basic-auth-file=/tmp/users/user-details.csv

Create the necessary roles and role bindings for these users:

1. ---
2. kind: Role
3. apiVersion: rbac.authorization.k8s.io/v1
4. metadata:
5. namespace: default
6. name: pod-reader
7. rules:
8. - apiGroups: ["" ] # "" indicates the core API group
9. resources: ["pods"]
10. verbs: ["get", "watch", "list"]
- 11.
12. ---
13. # This role binding allows "jane" to read pods in the "default" namespace.
14. kind: RoleBinding
15. apiVersion: rbac.authorization.k8s.io/v1
16. metadata:
17. name: read-pods
18. namespace: default
19. subjects:
20. - kind: User
21. name: user1 # Name is case sensitive
22. apiGroup: rbac.authorization.k8s.io
23. roleRef:
24. kind: Role #this must be Role or ClusterRole
25. name: pod-reader # this must match the name of the Role or ClusterRole  
you wish to bind to
26. apiGroup: rbac.authorization.k8s.io

Once created, you may authenticate into the kube-api server using the users credentials

```
curl -v -k https://localhost:6443/api/v1/pods -u "user1:password123"
```

Presenter: Hello and welcome to this lecture.

In this lecture, we look at kubeconfigs in Kubernetes.

So far, we have seen how to generate a certificate for a user.

We have seen how a client uses the certificate file and key to query the Kubernetes REST API, for a list of pods using curl.

In this case, my cluster is called my-kube-playground.

So, send a curl request

to the address of the kube API server,

while passing in the pair of files,

along with the CA certificate,

as options.

This is then validated by the API server to authenticate the user.

Now, how do you do that

while using the kubectl command?

You can specify the same information using the options,

server, client key, client certificate

and certificate authority with the kubectl utility.

Obviously, typing those in every time is a tedious task.

So, you move these information

to a configuration file called as kubeconfig,

and then specify this file

as the kubeconfig option in your command.

By default, the kubectl tool looks for a file

named config under a directory, .kube,

under the user's home directory.

So, if you create the kubeconfig file there,

you don't have to specify the path

to the file explicitly in the kubectl command.

That's the reason you haven't been specifying any options for your kubectl commands so far.

The kubeconfig file is in a specific format.

Let's take a look at that.

The config file has three sections,

clusters, users and contexts.

Clusters are the various Kubernetes clusters

that you need access to.

Say you have multiple clusters for development environment

or testing environment, or prod,

or for different organizations,

or on different cloud providers, et cetera.

All those go there.

Users are the user accounts with which you have access to these clusters.

For example, the admin user, a dev user, a prod user, et cetera.

These users may have different privileges on different clusters.

Finally, contexts marry these together.

Contexts define which user account will be used to access which cluster.

For example, you could create a context named admin@production that will use the admin account to access a production cluster.

Or, I may want to access the cluster I have set up on Google with the dev user's credentials, to test deploying the application I built.

Remember, you're not creating any new users, or configuring any kind of user access or authorization in the cluster with this process.

You're using existing users with their existing privileges,

and defining what user you're going to use to access what cluster.

That way, you don't have to specify the user certificates and server address in each and every kubectl command you run.

So, how does it fit into our example?

The server specification in our command goes into the clusters section.

The admin user's keys and certificates goes into the users section.

You then create a context that specifies to use the my-kube-admin user to access the my-kube-playground cluster.

Let's look at a real kubeconfig file now.

The kubeconfig file is in a YAML format.

It has API version set to V1.

The kind is config, and then it has three sections, as we discussed.

One for clusters, one for contexts, and one for users.

Each of these is in an array format.

That way, you can specify multiple clusters, users or contexts within the same file.

Under clusters, we add a new item for our kube-playground cluster.

We name it my-kube-playground, and specify the server address under the server field.

It also requires the certificate

of the certificate authority.

We can then add an entry into the users section to specify details of my kube admin user.

Provide the location of the client certificate and key pair, so we have now defined the cluster and the user to access the cluster.

Next, we create an entry under the contexts section to link the two together.

We will name the context my-kube-admin@my-kube-playground.

We will then specify the same name we used for cluster and user.

Follow the same procedure to add all the clusters you daily access.

The user credentials you use to access them, as well as the contexts.

Once the file is ready, remember that you don't have to create any object like you usually do for other Kubernetes objects.

The file is left as-is, and is read by the kubectl command, and the required values are used.

Now, how does kubectl know which context to choose from?

We've defined three contexts here.

Which one should it start with?

You can specify the default context to use by adding a field current-context to the kubeconfig file.

Specify the name of the context to use.

In this case, kubectl will always use the context dev-user@google to access the Google clusters, using the dev user's credentials.

There are command line options available within kubectl to view and modify the kubeconfig files.

To view the current file being used, run the `kubectl config view` command.

It lists the clusters, contexts, and users, as well as the current context that is set.

As we discussed earlier, if you do not specify which kubeconfig file to use, it ends up using the default file located in the folder .kube, in the user's home directory.

Alternatively, you can specify a kubeconfig file by passing the kubeconfig option in the command line, like this.

We will move our custom config to the home directory, so this becomes our default config file.

So, how do you update your current context?

So, you've been using my-kube-admin user to access my-kube-playground.

How do you change the context to use prod user to access the production cluster?



Run the `kubect! config use-context` command to change the current context

to the prod user at production context.

This can be seen in the current context field in the file.

So yes, the changes made by `kubect! config` command actually reflects in the file.

You can make other changes in the file,

update or delete items in it,

using other variations of the `kubect! config` command.

Check them out when you get time.

What about namespaces?

For example, each cluster may be configured with multiple namespaces within it.

Can you configure a context to switch

to a particular namespace? Yes.

The `context` section in the `kubeconfig` file can take additional field called `namespace`,

where you can specify a particular namespace.

This way, when you switch to that context,

you will automatically be in a specific namespace.

Finally, a word on certificates.

You have seen path to certificate files mentioned in `kubeconfig` like this.

Well, it's better to use the full path,

like this, but remember, there's also another way to specify the certificate credentials.

Let's look at the first one.

For instance, where we configure the path

to the certificate authority.

We have the contents of the `ca.crt` file on the right.

Instead of using certificate authority field

and the path to the file,

you may optionally use the `certificate authority data` field

and provide the contents of the certificate itself.

But, not the file as-is.

Convert the contents to a base64 encoded format

and then pass that in.

Similarly, if you see a file

with the certificate's data in the encoded format,

use the `base64 decode` option to decode the certificate.

Well, that's it for this lecture.

Head over to the practice exercises section,

and practice working with `kubeconfig` files

and troubleshooting issues.

-: Before we head into authorization,

it is necessary to understand

about API groups in Kubernetes.

But first, what is the Kubernetes API?

We learned about the Kube API server.

Whatever operations we have done so far with the cluster we have been interacting with the API server one way or the other, either through the Kube control utility or directly via REST.

Say we want to check the version.

We can access the API server at the master nodes address followed by the port, which is 6443 by default and the API version. It returns the version.

Similarly, to get the list of pods you would access the URL API slash V1 slash pods.

Our focus in this lecture is about these API pods: the version and the API.

The Kubernetes API is grouped into multiple such groups based on their purpose such as one for APIs, one for health one for metrics and logs, et cetera.

The version API is for viewing the version of the cluster, as we just saw. The metrics and health API are used to monitor the health of the cluster.

The logs are used for integrating with third party logging applications.

In this video, we will focus on the APIs responsible for the cluster functionality.

These APIs are categorized into two: the core group and the named group.

The core group is where all core functionality exists such as namespaces, pods, replication controllers, events and points, notes, bindings, persistent volumes, persistent volume claims, conflict maps, secrets, services, et cetera.

The named group APIs are more organized.

And going forward, all the newer features are going to be made available through these named groups.

It has groups under it for apps, extensions, networking, storage, authentication, authorization, et cetera.

Shown here are just a few.

Within apps, you have deployments, replica sets, Stateful sets.

Within networking, you have network policies.

Certificates have these certificate signing requests that we talked about earlier in this section.

So the ones at the top are API groups, and the ones at the bottom are resources in those groups.

Each resource in this has a set of actions associated with them.

Things that you can do with these resources,

such as list the deployments, get information about one of these deployments, create a deployment, delete a deployment, update a deployment, watch a deployment, et cetera. These are known as verbs.

The Kubernetes API reference page can tell you what the API group is for each object. Select an object and the first section in the documentation page shows its group details. V1 core is just V1.

You can also view these on your Kubernetes cluster.

Access your Kube API server at port 6443 without any path and it will list you the available API groups.

And then within the named API groups, it returns all the supported resource groups. A quick note on accessing the cluster API like that.

If you were to access the API directly through curl as shown here, then you will not be allowed access except for certain APIs like version, as you have not specified any authentication mechanisms.

So you have to authenticate to the API using your certified files by passing them in the command line like this.

An alternate option is to start a Kube control proxy client.

The Kube control proxy command launches a proxy service locally on port 8001 and uses credentials and certificates from your Kube config file to access the cluster. That way you don't have to specify those in the curl command.

Now you can access the Kube control proxy service at port 8001 and the proxy will use the credentials from Kube config file to forward your request to the Kube API server. This will list all available APIs at root.

So here are two terms that kind of sound the same: the Kube proxy and Kube control proxy. Well, they're not the same.

We discussed about Kube proxy earlier in this course. It is used to enable connectivity between parts and services across different nodes in the cluster.

We discuss about Kube proxy in much more detail later in this course.

Whereas Kube control proxy is an HTTP proxy service created by Kube control utility to access the Kube API server.

So what to take away from this.

All resources in Kubernetes are grouped into different API groups.

At the top level,

you have core API group and named API group.  
Under the named API group, you have one for each section.  
Under these API groups, you have the different resources  
and each resource has a set  
of associated actions known as verbs.

In the next section on authorization,  
we can see how we use these to allow  
or deny access to users.

Well, that's it for this lecture.

I will see you in the next.

-: So far, we talked about authentication.

We saw how someone can gain access to a cluster.

We saw different ways that someone, a human or a machine  
can get access to the cluster.

Once they gain access, what can they do?

That's what authorization defines.

First of all, why do you need authorization in your cluster?

As an administrator of the cluster,  
we were able to perform all sorts of operations in it.

Such as viewing various objects like pods,  
and notes, and deployments,  
creating or deleting objects such  
as adding or deleting pods or even notes in the cluster.

As an admin, we are able to perform any operation.

But soon, we will have others accessing the cluster as well  
such as the other administrators, developers, testers,  
or other applications like monitoring applications  
or continuous delivery applications like Jenkins, et cetera.

So, we will be creating accounts  
for them to access the cluster by creating user names  
and passwords, or tokens, or signed TL certificates,  
or service accounts as we saw in the previous lectures.

But we don't want all  
of them to have the same level of access as us.

For example, we don't want the developers to have access to  
modify our cluster configuration, like adding  
or deleting notes or the storage or networking  
configurations.

We can allow them to view, but not modify.

But they could have access to deploying applications.

The same goes with service accounts.

We only want to provide the external application  
the minimum level of access to perform its required  
operations.

When we share our cluster between  
different organizations or teams, by logically  
partitioning it using name spaces, we want to

restrict access to the users to their name spaces alone.  
That is what authorization can help you within the cluster.  
There are different authorization mechanisms supported by Kubernetes, such as node authorization, attribute based authorization, role-based authorization, and Webhook.

Let's just go through these, now.

We know that the Kube API server is accessed by users like customer management purposes, as well as the kubelets on nodes within the cluster for management purposes within the cluster.

The kubelet accesses the API server to read information about services and points, nodes, and pods.

The kubelet also reports to the Kube API server with information about the node, such as its status.

These requests are handled by a special authorizer known as the node authorizer.

In the earlier lectures, when we discussed about certificates, we discussed that the kubelets should be part of the system nodes group and have a name prefixed with system node.

So any request coming from a user with the name system node and part of the system nodes group is authorized by the node authorizer and are granted these privileges, the privileges required for a kubelet.

So, that's access within the cluster.

Let's talk about external access to the API.

For instance, a user attribute based authorization is where you associate a user or a group of users with a set of permissions.

In this case,

we say the dev user can view, create and delete pods.

You do this, by creating a policy file with a set of policies defined in adjacent format.

This way, you pass this file into the API server.

Similarly, we create a policy definition file for each user or group in this file.

Now, every time you need to add or make a change in the security, you must edit this policy file manually and restart the Kube API server.

As such, the attributes based access control configurations are difficult to manage.

We will look at role based access controls, next

Role based access controls make these much easier.

With role based access controls,

instead of directly associating a user

or a group with a set of permissions, we define a role.

In this case, for developers, we create a role with the set of permissions required for developers. Then, we associate all the developers to that role. Similarly, create a role for security users with the right set of permissions required for them. Then, associate the user to that role, going forward. Whenever a change needs to be made to the user's access we simply modify the role and it reflects on all developers immediately. Role-based access controls provide a more standard approach to managing access within the Kubernetes cluster. We will look at role-based access controls in much more detail in the next lecture. For now, let's proceed with the other authorization mechanisms. Now, what if you want to outsource all the authorization mechanisms?

Say you want to manage a authorization externally and not through the built in mechanisms that we just discussed.

For instance, open policy agent is a third party tool that helps with admission control and authorization.

You can have Kubernetes make an API call to the open policy agent with the information about the user and his access requirements and have the open policy agent decide if the user should be permitted or not.

Based on that response, the user is granted access.

Now, there are two more modes in addition to what we just saw.

Always allow and always deny.

As the name states always allow, allows all requests without performing any authorization checks.

Always deny, denies all requests.

So, where do you configure these modes?

Which of them are active by default?

Can you have more than one at a time?

How does authorization work if you do have multiple ones configured?

The modes are set using the authorization mode option on the Kube API server.

If you don't specify this option, it is set to always allow by default.

You may provide a comma separated list of multiple modes that you wish to use.

In this case, I want to set it to node , Rback and Webhook.

When you have multiple modes configured, your request is authorized using each one in the order it is specified.

For example, when a user sends a request it's first handled by the node authorizer. The node authorizer handles only node requests. So, it denies the request. Whenever a module denies a request, it is forwarded to the next one in the chain. The role-based access control module performs its checks and grants the user permission. Authorization is complete and user is given access to the requested object. So, every time a module denies the request, it goes to the next one in the chain and as soon as a module approves the request, no more checks are done and the user is granted permission. Well, that's it for this lecture. In the upcoming lectures we will discuss more about role-based access controls.

-: In this lecture, we'll look at role-based access controls in much more detail. So how do we create a role? We do that by creating a role object. So we create a role definition file with the API version set to R back dot authorization dot K etes dot io slash V1, and kind set to role. We name the role developer as we are creating this role for developers, and then we specify rules. Each rule has three sections, API groups, resources, and verbs. The same things that we talked about in one of the previous lectures. For core group, you can leave the API group section as blank. For any other group, you specify the group name. The resources that we want to give developers access to are pods. The actions that they can take are list, get, create, and delete. Similarly, to allow the developers to create config maps, we add another rule to create config map. You can add multiple rules for a single role like this. Create the role using the kube control create role command. The next step is to link the user to that role. For this, we create another object called role binding. The role binding object links a user object to a role. We will name it dev user to developer binding. The kind is role binding.

It has two sections.

The subject is where we specify the user details.

The role ref section is where we provide the details of the role we created.

Create the role binding

using the kube control create command.

Also note that the roles and role bindings fall under the scope of name spaces.

So here the dev user gets access to pods and config maps within the default name space.

If you want to limit the dev user's access within a different name space, then specify the name space within the metadata of the definition file

while creating them.

To view the created roles, run the kube control get roles command.

To list role bindings, run the kube control get roll bindings command.

To view more details about the role, run the kube control describe role developer command.

Here you see the details about the resources and permissions for each resource.

Similarly, to view details about role bindings, run the kube control describe role bindings command.

Here you can see details about an existing role binding.

What if you being a user would like to see if you have access to a particular resource in the cluster?

You can use the kube control auth can I command and check if you can say, create deployments, or say, delete notes.

If you are an administrator, then you can even impersonate another user to check their permission.

For instance, say you were tasked to create a necessary set of permissions for a user to perform a set of operations and you did that,

but you would like to test if what you did is working.

You don't have to authenticate as the user to test it.

Instead, you can use the same command with the as user option like this.

Since we did not grant the developer permissions to create deployments, it returns no.

The dev user has access to creating pods though.

You can also specify the name space in the command like this.

The dev user does not have permission to create a pod in the test name space.

A quick note on resource names.



We just saw how you can provide access to users for resources like pods within the name space. You can go one level down and allow access to specific resources alone. For example, say you have five pods in name space, you wanna give access to a user to pods, but not all pods. You can restrict access to the blue and orange pod alone by adding a resource names field to the role. Well, that's it for this lecture. Head over to the practice exercises section and practice working with role-based access controls.

Instructor: We discussed about roles and role bindings in the previous lecture. In this lecture, we will talk about cluster roles and cluster role bindings. When we talked about roles and role bindings, we said that roles and role bindings are name spaced, meaning they're created within name spaces. If you don't specify a name space, they're created in the default name space and control access within that name space alone. In one of the previous lectures, we discussed about name spaces and how it helps in grouping or isolating resources like pods, deployments and services. But what about other resources like nodes? Can you group or isolate nodes within the name space? Like can you say node 01 is part of the dev name space? Now, those are cluster wide or clusters scoped resources. They cannot be associated to any particular name space, so the resources are categorized as either name spaced or cluster scoped. Now, we have seen a lot of name spaced resources throughout this course like pods and replica sets and jobs, deployments, services, secrets. And in the last lecture, we saw two new, roles and role bindings. These resources are created in the name space you specify when you created them. If you don't specify a name space, they're created in the default name space. To view them or delete them or update them, you always specify the right name space. The cluster scoped resources are those where you don't specify a name space when you create them like nodes, persistent volumes, the cluster roles and cluster role bindings

that we're going to look at in this lecture.  
Certificate signing request we saw earlier  
and name spaced objects themselves are  
of course, not name spaced.  
Note that, this is not a comprehensive list of resources.  
To see a full list of name spaced  
and non-name spaced resources,  
run the kube control API resources command  
with the name spaced option set.  
In the previous lecture, we saw how to authorize a user  
to name space resources.  
We use roles and role bindings for that,  
but how do we authorize users  
to cluster via resources like nodes or persistent volumes?  
That is where you use cluster roles  
and cluster role bindings.

Cluster roles are just like roles except they are  
for cluster scoped resources.  
For example, a cluster admin role can be created  
to provide a cluster administrator permissions  
to view, create, or delete nodes in a cluster.  
Similarly, a storage administrator role can be created  
to authorize a storage admin to create persistent volumes.  
Create a cluster role definition file  
with the kind cluster role  
and specify the rules as we did before.  
In this case, the resources are nodes,  
then create the cluster role.  
The next step is to link the user to that cluster role.  
For this, we create another object  
called cluster role binding.  
The role binding object links the user to the role,  
we will name it cluster admin role binding.  
The kind is cluster role binding.  
Under subjects, we specify the user details,  
cluster admin user, in this case.  
The role ref section is where we provide the details  
of the cluster role we created.  
Create the role binding using  
the kube control create command.  
One thing to note before I let you go,  
we said that cluster roles and bindings are used  
for cluster scoped resources, but that is not a hard rule.  
You can create a cluster role  
for name spaced resources as well.  
When you do that, the user will have access  
to these resources across all name spaces.  
Earlier, when we created a role to authorize a user  
to access pods, the user had access

to the pods in a particular name space alone.

With cluster roles, when you authorize a user to access the pods, the user gets access to all pods across the cluster.

Kubernetes creates a number of cluster roles by default when the cluster is first setup.

We will explore those in the practice tests that are coming up. Good luck.

### admission controllers.

So we've been running commands from our command line using the kubectl utility to perform various kinds of operations on our Kubernetes cluster.

And we know every time we send a request, say to create a pod, the request goes to the API server and then the pod is created and the information is finally persisted in the etcd database.

When the request hits the API server, we've learned that it goes through an authentication process and this is usually done through certificates.

If the request was sent through kubectl, we know the kubeconfig file has the certificates configured and the authentication process is responsible for identifying the user who send the request and making sure the user is valid.

And then the request goes through an authorization process, and this is when we check if the user has permission to perform that operation.

And we have learned that this is achieved through role-based access controls.

So in this case if the user was assigned this particular role of a developer the user is allowed to list, get, create, update or delete pods.

And so if the request that came in matched any of these conditions, in this case it does as the request is to create a pod.

It is allowed to go through, otherwise it's rejected.

So that's authorization with role based access control.

Now with role based access control, you could place in different kinds of restrictions, such as to allow or deny those with a particular role to say create list or delete different kinds of objects like pods, deployments or services.

We could even restrict access to specific resource names, such as say a developer can only work on pods named blue or orange, or restrict access within a namespace alone.

Now, as you can see, most of these rules that you can create with role-based access control is at the Kubernetes API level, and what user is allowed access to what kind of API operations.

And it does not go beyond that.

But, what if you want to do more than just define what kind of access a user has to an object?

For example, when a pod creation request comes in you'd like to review the configuration file and look at the image name and say that you do not want to allow images from a public dock hub registry.

Only allow images from a specific internal registry.

Or to enforce that we must never use the latest tag for any images.

Or say, for example you'd like to say, if the container is running as the root user, then you do not want to allow that request.

Or allow certain capabilities only, or to enforce that the metadata always contains labels.

So, these are some of the things that you can't achieve with the existing role-based access controls, and that is where admission controllers comes in.

Admission controllers help us implement better security measures to enforce how a cluster is used.

Apart from simply validating configuration, admission controllers can do a lot more, such as change the request itself or perform additional operations before the pod gets created.

We will go over some examples in the upcoming slides.

There are a number of admission controllers that come prebuilt with Kubernetes, such as always pull images that ensures that every time a pod is created the images are always pulled.

The default storage class admission controller that observes the creation of PVCs and automatically adds a default storage class to them if one is not specified.

The event rate limit admission controller can help set a limit on the request with the API server can handle at a time, to prevent the API server from flooding with requests.

The namespace exists, admission controller rejects requests to namespaces that do not exist,

and there are many more admission controllers available.

So let's take that as an example and look at it in a bit more detail.

The namespace exists admission controller.

Say we want to create a pod in a namespace

called blue that doesn't exist.

If I run this command, it would throw an error

that says the namespace blue is not found.

What's happening here is that my request gets authenticated then authorized, and it then goes through the admission controllers.

The namespace exists that mission controller handles the request and checks if the blue space is available.

If it is not, the request is rejected.

The namespace exists is a built in admission controller that is enabled by default.

There's another admission controller that is not enabled by default, and that is called as the

Namespace Auto Provision Admission controller.

This will automatically create the namespace, if it does not exist.

We will see how it can be enabled in a minute.

First, to see a list of admission controllers enabled by default run the kube API server/H command and grip for enable admission plugins.

Now here you'll see a list of admission controllers that are enabled by default the ones that are highlighted in green.

Note that if you're running this in a Kube ADM based setup, then you must run this command within the Kube APA Server Control Plane Pod using the kubectl exec command first like this.

To add an admission controller update the enable admission plugins flag on the Kube APA server service to add the new admission controller.

So if you're in a kube ADM based setup then update the flag within the kube APA server manifest file as shown here on the right.

So the one on the left is

if you are updating the kube APA server service

and the one on the right is

if the APIs server is running as a pod

in a kube ADM based setup.

Similarly to disabled admission controller plugins, you could use the disabled admission plugins flag.

Once updated the next time we run the command to provision a pod in a namespace that does not exist yet, the request goes through authentication, then authorization, and then the namespace auto probation controller at which point it realizes that the namespace doesn't exist.

So it creates the namespace automatically

and the request goes through successfully to create the pod.

If you list the namespaces now,  
you'll see that the blue namespace is automatically created.  
So that's one example of how an admission controller works.  
It cannot only validate and reject requests from users.  
It can also perform operations in the backend  
or change the request itself.  
Note that the namespace auto provision  
and the namespace exists admission controllers  
are deprecated and is now replaced by the namespace  
lifecycle admission controller.  
The namespace lifecycle admission controller will make sure  
that request to a non-existent namespace is rejected  
and that the default namespaces such as  
default kube system and kube public cannot be deleted.  
Well, that's it for now. Head over to the labs and practice  
are working with admission controllers and I will see you  
in the next lecture.

-: In this lecture, we will take a closer look  
at the different types of admission controllers  
and how we can configure our own admission controller.  
We looked at the namespace exists  
or namespace lifecycle admission controller.  
It can help validate if a name space already exists  
and reject the request if it doesn't exist.  
This is known as a validating admission controller.  
Let's look at another type  
of admission controller plugin  
named as the default StorageClass plugin.  
This is a plugin that is enabled by default.  
Say for example,  
you're submitting a request to create a PVC.  
The request goes through authentication authorization  
and finally, the admission controller.  
The default StorageClass admission controller  
will watch for request to create a PVC  
and check if it has a StorageClass mentioned in it.  
If not, which is true in our case,  
it'll modify your request to add the default StorageClass  
to your request.  
This could be whatever StorageClass is configured  
as the default StorageClass in your cluster.  
So when the PVC is created and you inspect it,  
you'll see that a StorageClass default is added to it  
even though you hadn't specified it during the creation.  
So this type of admission controller  
is known as a mutating admission controller.  
It can change or mutate the object itself

before it is created.

So those are two types of admission controllers.

Mutating admission controllers are those that can change the request, and validating admission controllers are those that can validate the request and allow or deny it.

And there may be admission controllers that can do both.

That can mutate a request as well as validate a request.

Now, generally, mutating admission controllers are invoked first

followed by validating admission controllers.

This is so that any change made by the mutating admission controller can be considered during the validation process.

And this example,

the namespace auto provisioning admission controller which is a mutating admission controller, is run first followed by the validating controller namespace exists.

If it was run the other way, then the namespace exists admission controller would always reject the request for a namespace that does not exist and the namespace auto provisioning controller would never be invoked to create the missing namespace.

And when a request goes through these admission controllers, if any admission controller rejects the request, the request is rejected and an error message is shown to the user.

Now, these are all built-in admission controllers that are part of the Kubernetes source code and are compiled and shipped with Kubernetes.

Now, what if we want our own admission controller with our own mutations and validations that has our own logic?

To support external admission controllers, there are two special admission controllers available: mutating admission webhook, and then validating admission webhook.

And this is what we will look at next.

We can configure these webhooks to point to a server that's hosted either within the Kubernetes cluster or outside it,

and our server will have our own admission webhook service running with our own code and logic.

After our request goes through, all the built-in admission controllers it hits the webhook that's configured.

We will see how to configure that in a bit.  
And then once it hits the webhook,  
it makes a call to the admission webhook server  
bypassing in an admission review object in adjacent format.

This object has all the details  
about the request such as the user that made the request,  
and the type of operation the user is trying to perform,  
and on what object and the details  
about the object itself.

On receiving the request,  
the admission webhook server responds  
with an admission review object  
with a result of whether the request is allowed or not.  
If the allowed field in the response is set to true  
then the request is allowed,  
and if it's set to false it is rejected.

So how do we set this up?

First, we must deploy our admission webhook server  
which will have our own logic,  
and then we configure the webhook on Kubernetes  
by creating a webhook configuration object.

So let's take a look at each of these steps next.

So the first step is to deploy our own webhook server.

Now, this could be an API server  
that could be built on any platform.

An example, [code of a webhook server written in Go](#),  
now can be found here in the Kubernetes documentation pages.  
It's written in the Go programming language.

You could develop your own server in other languages  
as well if required.

The only requirement is that it must accept the mutate  
and validate APIs and respond with the JSON object  
that the web server expects.

So here's a pseudo code of a sample webhook server  
written in Python.

There are two calls, a validate call and a mutate call.  
The validate call receives the validation webhook request  
and in this example compares the name of the object  
and the name of the user who sent the request  
and rejects the request if it's the same name.

Well, just a simple example use case to show what we can do  
with the request that come in.

And if you look down, we'll see the mutating webhook  
which gets the username and response  
with a JSON Patch operation of adding the username  
as a label to any request that was raised by anyone.

So if you take a closer look  
at this particular piece of code,  
a patch object is a list of patch operations



with each operation being add, remove, replace, move, copy or test,  
and we then specify the path  
within the JSON object that needs to be targeted for change.  
In this case, it is /metadata/label/users,  
and then the value that needs to be added  
if it is an add operation.

So we get the username from the request.  
So that's gonna be the value of that particular label.  
This is then sent as a base 64 encoded object  
as part of the response.

On a side note, from an exam point of view,  
you will not be asked to develop any code like this,  
so don't worry if you don't fully understand  
this piece of code.

All you need to take away from this  
is that the admission webhook server  
is a server that you deploy that contains the logic  
or the code to permit or reject a request  
and it must be able to receive  
and respond with the appropriate responses  
that the webhook expects.

So this is just a simple example to show  
what kind of things that you can do  
or what kind of things that you can code  
and implement in the webhook server that you deploy.  
Okay.

So moving on, once we have developed our own webhook server  
the next step is to host it.

So we either run it as a server somewhere  
or containerize it and deploy it  
within Kubernetes cluster itself as a deployment.

If deployed as a deployment in a Kubernetes cluster,  
then it needs a service for it to be accessed.

So we have a service named webhook service as well.

The next step is to configure our cluster  
to reach out to the service  
and validate or mutate the requests.

For this,  
we create a validating webhook configuration object.

So we start with the API version,  
kind, metadata and webhooks section.

API version is admissionregistration.k8s.io/v1,  
kind is validating webhook configuration.

If we are configuring a mutating webhook,  
this would be a mutating webhook configuration.

We then give it a name.

Now under webhooks, we configure the different webhooks.

So a webhook has a name,

a client config and a set of rules.

So the name we set it to podpolicy.example.com.

The client config is where we configure the location of our admission webhook server.

If we deploy this server externally on our own that is not a part of a deployment in Kubernetes cluster, then we can simply provide a URL path to that server like this.

Now instead, if we deployed the server as another service on our own cluster as we see here as it is on the left, then we can use the service configuration and provide the namespace, and name of the service, which in our case is webhook-service.

Now, of course, the communication between the API server and the webhook server has to be over TLS.

So a certificate bundle should be configured.

So the server has to be configured with a pair of certificates.

Then a CA bundle is to be created and passed into this client config SSEA bundle.

Next, we must specify when to call our API server.

We could specify rules to configure exactly when we want our web hook server to be called for validation.

Now, we might not want to do that for all of the calls: for example, we may only want it to be called while creating pods, or deleting pods, or creating deployments, et cetera.

Now, whatever that may be, it can be added as a rule under the rules section using API groups, API versions, operation types and resources.

In this example, we are only going to call this webhook configuration when calls are made to create pods.

And that should be it.

Once this object is created, every time we create a pod a call would be made to the webhook service, and depending on the response, it would be allowed or rejected.

Well, this is for this lecture.

Head over to the labs and practice working with webhooks and I will see you in the next lecture.

Autoscroll

Course content

Overview

Q&A Questions and answers

Notes

Announcements

## Reviews

### Learning tools

Instructor: We will now discuss about API versions.  
We talked about APIs and API groups, resources and verbs.  
Now we are going to talk about API versions.  
We know that everything under APIs are the API groups such as apps, extensions, networking, et cetera.  
And each API group has different versions.  
Now, what we see here is the V1 version.  
When an API group is at V1 that means it is a GA stable version generally available stable version.  
The API group may have other versions such as Beta or Alpha as V1 Beta 1 or V1 Alpha 1 respectively.  
So what are these different versions?  
What do these each of these versions mean?  
Now Alpha is when an API is first developed and merged to the Kubernetes code base and becomes part of the Kubernetes released for the very first time.  
At this stage, the API version has Alpha in its name indicating that it is an Alpha release.  
For example, it could be V1 Alpha 1 or V1 Alpha two, et cetera.  
This API group is not enabled by default.  
For example, at the time of this recording, the API group internal dot API server dot case dot io, which has the resource storage version is only available as V1 Alpha 1.  
This means that if you were to create this object now you will have to say in the amul definition file the API version as internal API server dot case dot io/ V1 Alpha 1.  
However, this Alpha version is not enabled by default.  
So if you try to create this object you will not be able to do that.  
We'll discuss about how to enable a particular API version specifically those in the Alpha version, a few minutes later in this video.  
Now, since this is an Alpha version it may lack end to end tests and may have bugs.  
As such, it is not very reliable.  
Also, there is no guarantee that this API will be available in the future releases.  
It may be dropped without any notice in future releases.  
As such, this is only really for expert users who are interested in testing and giving early feedback

for the API group.

As of this recording

the API groups in Alpha phase

are the internal API server dot case dot IO group.

Now, after some time the Alpha version

and once all the major bug are fixed

and it has end to end test, it moves to the Beta stage.

This is where the API version name gets Beta in name.

For example,

V1 Beta one or V1 Beta two, et cetera.

The API groups in Beta stage

are enabled by default

and they have end-to-end tests.

But since it is not GA,

it may still have minor bugs.

However, there is commitment

from the project that it may be moved to GA in the future.

This is for users who are interested

in Beta testing and providing feedback for the features.

As of now, the flow control group is in Beta phase,

as as you can see from this chart.

Now, after being in the Beta stage

for a few months and with a few releases

the API group moves to the GA or stable version.

Now, this is when the version number no longer

has Alpha or Beta in it.

Instead, it is just V1, and it is of course enabled

by default in the API group,

and it is part of conformance tests

and has all the tests written.

Now, since most bugs are fixed

in Alpha and Beta stages in the API group,

this API group is highly reliable

and it'll be supported and present

in many feature releases to come.

And of course, this is when it is available to all users.

Now, as you might have noticed

most of the API groups such as apps, authentication

authorization certificates, coordination, et cetera

are GA and stable now.

However, if you look at some of them,

they have multiple versions.

For example,

autoscaling has V1, V2, Beta 2, V2, Beta 1, et cetera.

The node API group has V1

V1 Beta 1 and V1 Alpha versions.

And so why is that and what does that mean?

So an API group can support multiple versions

at the same time.

For example, the apps group can have..

V1 Beta 1 and V1 Alpha 1

all at the same time.

This would mean

that you will be able to create the same object

using any of these versions in your YAML file.

So say I want to create an nginx dot yaml,

I could create it

with the API version set to V1 Alpha 1

or V1 Beta 1 or V1.

So any of these will work.

But even though there are multiple versions supported

at the same time,

only one can be the preferred or storage version.

So what is a preferred version?

Now, when you have multiple versions enabled

and you run the `kubectl get deployment` command

which version is the command going to query?

That's defined by the preferred version.

In this case, if V1 is set to the preferred version

then that is the command that it will query.

Or when you run the `kubectl explain` command

the version that it returns is the preferred API version.

Also, when multiple versions are present,

only one version can be the storage version.

This means if any object is created with the API version set

to anything other than the storage version,

such as V1 Alpha 1 or V1 Beta 1,

then those will be converted to the storage version,

which is V1

before storing them into the etcd database.

So that's what the preferred and storage versions mean.

Preferred is the default version used

when you retrieve information through the API

using `kubectl get` command or something like that.

And the storage version is the version

in which an object is stored

in etcd respective of the version you have used

in the yaml file while creating the object.

Remember, only one can be the preferred or storage version.

And even though the preferred

and storage versions are usually the same

they can be different.

They don't necessarily have to be the same.

So how do you identify the preferred version

if there are multiple APIs?

So the preferred version is listed when you list

that specific API.

For example, when I list the APIs available

for the batch API group,  
by going to the URL as shown here  
I see that it has V1 and V1 Beta 1 versions available  
which are the supported API versions  
of which the preferred version is V1 as is shown below.  
And as of now

it is not possible to see which is the storage version  
of a particular API through an API or a command.

The one way to find that out  
is by looking at the stored object  
in etcd itself by querying the etcd database directly.

Here's a sample command where we are querying  
the blue deployment object  
using the etcd kubectl utility using the get command.  
And in the output we see the API version  
to be apps slash V1.

So to enable or disable a specific version  
you must add it to the runtime config parameter  
of the Kube API server service.

For example, we talked about a number of Alpha APIs earlier  
and since these APIs are not enabled by default  
you will not be able to create objects for them.

So if you'd like to try those out  
what you could do is you could actually add those  
to the runtime config parameter  
and the APIs here are comma separated  
and you would like to enable  
and specify the APIs that you would like to enable.

And of course, once you do that  
you must remember to always restart the API server service.

And that goes to whenever you edit these options.

And once done  
you should be able to use that API  
and test it out.

Well, that's all for now  
and I will see you in the next lecture.

-: We will now discuss about API deprecations.

So we discussed  
that a single API group can support multiple versions  
at a time, but why do you need to support multiple versions  
and how many should you support?

When can you remove an older version that is  
no longer required?

This is answered by the API deprecation policy.

By looking at the answers to these questions  
we will also understand the rules  
in the API deprecation policy.

First, let's look at the why.

Let me give you a quick tour

of the life cycle of an API group.  
Let's for example say that we are planning to contribute to the Kubernetes project.  
So we create an API group called code cloud.com under which we have two resources called Cores and Webinar.  
This is just an example.  
So we develop it in-house and we test it and when we are ready to merge it to the Kubernetes project we raise a PR and hopefully it gets accepted and we release it as an alpha version.  
So we call it V1 Alpha one because it's the first alpha version of the V1 version you can create a course or a webinar object using yaul file like this and using the API version code cloud.com/v1 alpha one.  
Now, let's say for instance the webinar object didn't go well with the users and we decided to remove it.  
In the next Kubernetes release.  
can we just remove it from the V one alpha one version?  
No. That's where the first rule of API application policy comes into play.  
API elements may only be removed by incrementing the version of the API group meaning you can only remove the webinar element from the WE one alpha two version of the API group.  
It will continue to be present in the V one alpha one version of the element.  
So at this point, the resource in database is still at B one Alpha one but our API version has now changed to B one alpha two.  
So this is now going to be a problem.  
We will need to go back and change all API versions in our files from B one alpha one to V one alpha two which is why the new releases must support both V one alpha one and V one alpha two.  
But the preferred or storage version could be we one alpha two.  
This means that the users can use the same yamo files to create the resources, but internally it would be converted and stored as V one alpha two.  
So that brings us to the second rule of API deprecation policy.  
API objects must be able to round trip between API versions in a given release without information loss with the exception of whole risk resources that do not exist in some versions.  
If we create an API, an object in the V one alpha one version

and it is converted to V one alpha two  
and then back to V one alpha one, it should be the same  
as the original V one alpha one version.  
For example, we have a course object in B one alpha one.  
It has a spec field called type set two video.  
This is then converted  
to B one alpha two with B one alpha two  
we introduced a new field called Duration.  
This field was not there  
in we one alpha one where we now convert this back  
to B one Alpha one.  
It will now have the new field, but the original one didn't.  
So we must add an equivalent field  
in B one alpha one version so  
that the converted B one alpha one matches the  
original B one alpha one version.  
So continuing with our story  
we fixed some more bugs and are now ready for beta.  
Our first beta version called v1, BTA one is now ready  
and then after a few months we release the next beta version  
of v1 beta two.  
And finally we release our stable version, the GA version  
which we call as v1.  
So that's kind of how an API evolves over time.  
Starts with we want alpha one, then alpha two  
you required more alpha versions and then to beta versions  
beta one, beta two, and more beta versions if required.  
And finally we want  
now we don't have to have all the versions available  
at all times like this.  
Of course, we must deprecate  
and remove older versions as we release newer versions.  
So let's look at what are the rules  
and best practices around that.  
So let's say with the Kubernetes release of version X  
and we released the B one alpha one  
version of our API group.  
So X here could be Kubernetes version release 12.1 0.1  
for example.  
And that's the first time our package was included.  
So we made it version V one alpha one.  
Now since it is the only version it is  
the preferred and storage version.  
With the Kubernetes release of x plus one  
we released the V1 alpha two version of our API group.  
Now, since we are in the alpha phase  
we are not required to keep the older V1 alpha one version  
as part of this new release.  
And this is part of the rule



for a of the Kubernetes deprecation policy.  
This rule states that other  
than the most recent API versions in each track  
older API versions must be supported  
after they're announced deprecation for a duration  
of no less than GA 12 months  
or three releases for GA and nine months  
or three releases for beta and zero releases for alpha.  
So alpha versions need not be supported  
for any release, but beta and GA versions once  
released must be supported anywhere from nine to 12 months.

So that's why the release x  
plus one does not have V1 alpha one.  
So that may break things  
for those who had previously used the V1 alpha one  
as as we discussed earlier in this in this video.  
And so you must mention about this change  
in the release node of that Kubernetes version.  
Here's an example of one such mention  
of the removal of V1 alpha two configure API  
and the release node requests users to  
convert the V1 alpha two to V1 alpha three.  
In that example, we will talk  
about how to do just that a bit later in this video.  
So now in the x plus two version  
we release the first VI version  
of our API group is called V one, beta one.  
And now since the previous version was an alpha version  
which is V one alpha two, it is not required  
for that version to be part of the new release.  
And same as before  
the release notes must be updated to notify users to migrate  
from B one alpha two to V one beta one.  
Now that we are in beta going forward  
things are going to change because rule four A states  
that beta versions need to be supported  
for nine months or three releases, whichever is longer.  
So continuing to x plus three release.  
Now on x plus three, we have, we want beta two released.  
Now, as we discussed earlier from the beta version onwards  
it is required to have the previous beta release  
which is V1 beta one in this case as part of this release.  
So this release has both the new beta version as well  
as the previous beta version  
as part of this release.  
The V1 beta one version is now deprecated a note  
that we want beta one version is only deprecated  
but not removed.  
It's still part of this release

it's still going to be there  
for a couple of releases before it, it is actually removed.  
If you were to use the we want beta one API at this point  
it would display a deprecation warning.  
Now also note that at this time we want beta one  
is continuing to be the preferred version.  
So why is that?  
Why isn't the, v one beta two version  
the preferred and storage version  
because rule four B states that the preferred API version  
and the storage version for a given group may not advance  
until after a release has been made that supports  
both the new version and the previous version.  
So in this case  
the current release is the first release where  
both new and previous version  
that's beta one and beta two are supported.  
So we cannot change the preferred storage version yet.  
Instead, we must wait till the next release.  
The next release also has both versions except  
beta two now becomes the preferred storage version.  
But continuing on with version X plus five  
we now finally have the V one GA stable version.  
But along with that, we also have V one beta one  
and V one beta two versions.  
V one beta two is still the preferred or storage version  
because this is the first release that has  
the V one version, but it now becomes deprecated.  
So now V one beta one and V one beta two are deprecated.  
In the next release, we can remove V one beta one version  
because as per rule four A  
a beta version needs to be supported for three releases.  
beta one version was deprecated in release X plus three  
and it's been around for X plus four and X plus five.  
So that's three releases  
and it can now be removed.  
With x plus six V one can now  
be the preferred and storage version  
while beta two is going to stay around  
for another release to complete its three release support  
with X plus seven version, nothing changes  
as beta two continues to stay around  
for the last time before it will be removed.  
And in x plus eight version, beta two is removed  
and we are left with just the GA stable V one version.  
That's kind of how a version progresses.  
Now, let's say for example with the next release, we are now  
for the first time having the first version of V two  
so we decided to start working on a V2 version

and we now have the V2 Alpha one version available for the first time.

Now, if you look at the releases in the past every time we had a new version available, we deprecated we immediately deprecated the older version.

But now that we have the V two alpha one version ready can we deprecate the V1 version now?

No, because that's where rule three comes in.

And rule three says an API version in a given track may not be deprecated until a new API version, at least as stable is released. Meaning GA versions can deprecate beta or alpha versions but not the other way around.

An alpha version aren't deprecate a GA version in this case the V2 alpha one is an alpha version and V one is a GA version, so we cannot deprecate V one now.

The V two alpha one version needs to go through its life cycle of becoming a V two, alpha two then V two beta one, then v2 beta two, and then V two and and only can the V2 version deprecate the V1 version.

So that brings us to the final topic of this video, the cube converts command. Now, as we have been seeing when Kubernetes clusters are upgraded, we have new APIs being added and old ones being deprecated and removed. Now, when an old API is removed, it is important to note that you have to update your existing manifest files to the new version.

For example, say I have a yamo file with a deployment definition of v1 beta one and when Kubernetes is upgraded the beta one version of deployment is removed. So I would need to use the V1 version going forward. However, I may have a lot of yamo files in the old manifest in the old version, which has V1 beta one in it. So to convert my YAML files to newer version you may use the cube cuttle convert command and specify the old yamo file and the new output format.

For example, to convert this deployment definition file to the new version, run the cube convert command and specify the old file name followed by the new version which happens to be app slash B one.

And this will output the new version of the manifest definition in YAML format on screen.

Note that the cube convert command may not be available on your system by default and that's because it's a separate plugin.

So the Cube keke convert command is a special plugin that you have to install.

The installation instructions can be found in the Kubernetes documentation pages, along with the instructions for installing the kubelet utility. And so it's available at this link. We will go over this in the upcoming labs. Well, thanks for your time and I will see you in the next one.

Instructor: Let's look at custom resource definitions in Kubernetes.

Before we talk about custom resources, let's first talk about resources.

So here's a deployment definition file.

And when we create a deployment, Kubernetes creates a deployment and stores its information in the ETCD data store.

And we can create this deployment and then list the deployment and see the status.

And we can run the delete command to delete the deployment.

Now all of this is simply going to create, list, and modify the deployment object or resource in the ETCD data store.

But we know that when we create deployments it creates pods equal to the number of replicas defined in the deployment.

In this case, three.

So who or what is responsible for that?

And that's the job of a controller.

In this case, the deployment controller.

Now we don't have to create a controller because the deployment controller

is built in to the Kubernetes and it's already available.

So the controller is a process that runs in the background and its job is to continuously monitor the status of resources that it's supposed to manage.

In this case, the deployment that we created.

And when we create, update, or delete the deployment, it makes the necessary changes on the cluster to match what we have done.

In this case, when we create a deployment object, the controller creates a replica set which, in turn, creates as many pods as we have specified in the deployment definition file.

That's the job of a controller.

And here is how it looks in code.

So the deployment controller is written in Go and is part of the Kubernetes source code.

You don't have to understand Go

or understand this code for now.

I'm just showing this

so you can see what a controller actually looks like.

Okay, now we've learned about many resources throughout this course such as replica sets, deployments, jobs, cron jobs, stateful sets, namespace, et cetera. These are just few of the many resources available on the cluster.

And these have controllers that are responsible for watching the status of these objects and making the necessary changes on the cluster to maintain the state as expected.

Now let's do something fun.

Let's just like how we created the deployment to deploy multiple parts in a cluster, we would like to create, say, a flight ticket object to book a flight ticket.

I'm just picking this use case to show that this could be really anything, like literally anything.

And later towards the end of the video or end of this section, we will see some real life use cases.

Let's stick with this for now.

And we are first, we are going to see what we want and later, we will see how to achieve it.

So what I want to do is create an object called flight ticket.

We will say the API version is flights.com/v1.

The kind is FlightTicket. We'll name it my-flight-ticket.

The spec section will have some properties required to book the ticket such as a from airport, which we will set to Mumbai, and a to address, which we'll set to London, and also another property called number to specify the number of tickets that I want to book and that is two.

When I create this object,

I want to have a flight ticket resource created.

And when I list all flight tickets,

I want all flight tickets to be listed.

And when I delete a flight ticket,

I want the flight ticket booking to be deleted.

Now how we're going to do this, I'll explain in a bit.

For now, we're just discussing the what. What do we want?

So this is going to create or delete the flight ticket object in the ETCD data store.

But it's not actually going to book a flight ticket, is it?

We want this to actually go out

and book a flight ticket for real.  
Say for instance, there is this API available at, say,  
`book-flight.com/api`

that we can call to book a flight ticket.

So how do we call this API

whenever we create a flight ticket object  
to book a flight ticket?

So for that, we're going to need a controller.

So we'll create a flight ticket controller

and it's written in Go,

and we will watch for the creation or deletion  
of the flight ticket resource.

And when we create a resource,

it will contact the book flight ticket API

to book a flight ticket.

And when we delete the resource,

it would make an API call to delete that booking.

And it might look something like this in code.

So that's how it works at a high level.

So the flight ticket object that we created  
is a custom resource and the controller that we wrote  
to book the actual flight ticket by calling the API,  
that is called as a custom controller.

We have a custom resource and we have a custom controller.

Now let's see how we are going to achieve this.

If you tried to create a flight ticket resource now  
on your Kubernetes cluster,

you will see that it fails with the error message

that says there are no matches for the kind of flight ticket  
in version flights.com/v1.

Now this is because you can't simply create any resource  
that you want without configuring it in the Kubernetes API,  
without first telling Kubernetes that it should allow us  
to create a flight ticket object.

So that's what we have to do first.

You have to first define what that resource is  
that we want to create.

So for that, we need what is known as  
a custom resource definition or CRD.

We're going to use CRD to tell Kubernetes  
that we would like to create objects  
of kind FlightTicket going forward.

So the CRD is a similar object  
with API version, kind, metadata, and spec.

The API version is `apiextensions.k8s.io/v1`.

Kind is `CustomResourceDefinition`.

Metadata has name set to `flighttickets.flights.com`,  
let's say.

And under the spec section, we will specify a scope.

Now scope defines if the object is namespaced or not, and we know that there are namespaced and non-namespaced scopes in Kubernetes. For example, pods, replica sets, deployments are all scoped, whereas persistent volumes, nodes, and namespace itself are non-scoped objects.

So here we can define if this object is going to be namespace scoped or not.

We'll set it to namespaced for now.

Then we define groups.

Groups is the API group that we provide in the API version. It'll be flights.com.

And then we specify the names of the resource.

We define the kind, which is the kind that we used in the flight ticket definition file and we will set it to FlightTicket.

And we must also specify a singular and plural versions of names.

So the singular name here is just the flight ticket which is used to display the resource type in the output of the kubectl commands.

The plural is what is used by the Kubernetes API resource.

And if you run the kubectl API resources command, this is what is going to be shown here in the plural format.

So it's flight tickets there,

and, of course, you can provide a short version of the name.

Let's say FT.

This way we can just refer to the resource as FT when we run the kubectl commands.

Next we have versions.

Each resource can be configured with multiple versions as it passes through the various phases of its lifecycle.

For example, if it's a new resource type that we are going to introduce it, then we start at alpha or beta versions before making its way to version v1.

This is something that we discussed in the API versions lecture earlier in this course.

Let's name it version v1 for now.

And with multiple versions configured for the same resource, we must choose which ones are served through the API server and we also define which is the storage version.

If you have multiple versions, only one version can be marked as the storage version.

We discussed about what are storage versions in the lecture about APIs earlier in this course.

Next, we specify the schema.

The schema defines all the parameters that can be specified under the spec section when you create the object.

It defines what fields are supported and what type of value that fields supports, et cetera.

The schema uses OpenAPI V3 schema version.

We specify the different properties using an object type.

We have from, to, and number.

And from and to are string. Number is an integer.

We can also specify validations like the minimum or maximum value that can be specified under the number.

If the value is entered by the user does not fall within this range, the resource will not be created and will return an error message.

We can now create the custom resource definition by running the kubectl create command.

And once the custom resource definition is created, you can now create the flight ticket object and then get or delete it.

So that solves the first part of our problem, being able to create any kind of object type that we want on Kubernetes.

You can use custom resource definitions, or CRDs, to create any kind of resource you want on Kubernetes and specify a schema and add validations.

But it's only going to allow you to create these resources and store them in ETCD.

It's not actually going to do anything about these resources yet

because we don't yet have a controller for it and that's what we will look at next.

The second part is building a custom controller that can handle these resources that can watch when these resources are created and perform actions based on the resource specifications.

You see, without a controller, the custom resource that we created is just going to sit there.

It's just data locally in our ETCD data stores and does not actually do anything.

In the upcoming video, we'll see how to get started with creating such a controller.

-: In this lecture

we will look at developing Custom Controllers.

So to pick up from where we left off,

we have created a custom resource definition, and we are able to create our Flight Ticket objects.

And the data about the ticket is stored in ETCD.

Now what we need to do is monitor the status



of the objects in ETCD  
and perform actions such as making calls  
to the flight booking API to book, edit,  
or cancel flight tickets.

And that's why we need a Custom Controller.

So a controller is any process or code that runs in a loop  
and is continuously monitoring the Kubernetes cluster  
and listening to events of specific objects being changed,  
in this case, the flight ticket object.

And now we could do that in any way we can.

Of course you need to write some code.

So say I know Python well.

So I could write a code in Python  
that would query the Kubernetes API server  
for objects in the Kubernetes API.

And then I could watch for changes  
and make further calls to the API  
to make changes to the system.

However, developing a controller in Python  
may be challenging,  
as the calls made to the APIs may become expensive,  
and we will need to create our own queuing  
and caching mechanisms.

Developing in Go with the Kubernetes Go Client,  
provides support for other libraries  
like shared informers that provide caching  
and queuing mechanisms  
that can help build controllers easily  
and in the right way.

So how do you start building a Custom Controller?

So there's a GitHub repo named sample-controller.

So first clone this repo,  
and then we modify the controller,

.go with your custom code.

Then we build and run it.

So make sure you have the Go programming language  
installed on your machine.

Install it if it is and if it is isn't already installed.

And then clone the GitHub repo of the sample-controller.

We then see the into the sample controller directory,  
and then we customize the controller.go  
with our custom logic.

And we're not going to get into the details  
about the code itself at this point.

We will probably go over it at some other time  
where we will focus on advanced topics  
that would include like building controllers and operators.

For now, let's just assume

that we've customized the controller.go code.

And somewhere within the code,  
we're making a call to the flight booking API  
and to book flight tickets.  
And then once we do that, we then build the code,  
and then we run it.

And when we run it, we specify the kube config file  
that the controller can use to authenticate  
to the Kubernetes API.

And then the control process starts locally  
and it then watches for the creation  
of the flight ticket objects  
and makes the necessary pulse.

Once your controller is ready,  
you can decide on how to distribute it.  
Now you don't want to build and run it each time.

So you may package the Custom Controller  
in a Docker image and then choose  
to run it inside your Kubernetes cluster  
as a pod or a deployment.

That's a high level overview  
of building a Custom Controller.

Now in the exam, I don't expect them to ask a question  
about building a Custom Controller  
as it requires more coding knowledge.

So I don't anticipate a question on this.

However, there may be questions  
to build custom resource definitions  
and work with customer resource definition files,  
right, or to work with existing controllers  
that are already there.

It's good to know.

All right, so that's all for now.

And in the next video, we will talk about operators.

-: In this section, we will talk about Helm.

Now, Kubernetes is awesome  
at managing complex infrastructures.

We humans tend to struggle with complexity  
though applications we deploy  
into our Kubernetes cluster can become very complicated.

A typical app is usually made up  
of a collection of objects  
that need to interconnect to make everything work.

For example, even a relatively simple WordPress  
site might need the following,  
a deployment to declare the pods that you wanna run  
like with MySQL database server or the web servers.

A persistent volume to store the database,

a persistent volume claim,  
a service to expose the web server running in a pod,  
and a secret to store the admin password  
and maybe even more  
if you want extra stuff like periodic backups and so on.  
And for every object, we might need a separate YAML file.  
Then we need to kube coddle apply to every YAML file.  
Now, this can be a tedious task,  
but it's not the end of the world.  
Now imagine we download these YAML files from the internet.  
We're not happy with the defaults,  
so we start changing stuff.  
Persistent volumes are 20GB,  
but we know our website will need much more storage  
and we go to the YAML files where persistent volumes  
and claims are declared.  
We change 20 to 100 and more stuff we wanna change.  
We'll have to open up every YAML file  
and edit each one according to our needs.  
Well, not bad enough yet.  
Now imagine two months go by.  
We now have to upgrade some components in our app,  
and back to editing multiple YAML files with great care.  
So we don't change the wrong thing in the wrong place.  
We need to delete the app.  
We'll need to remember each object that belongs  
to our app and delete them all one by one.  
Now you might be thinking,  
hey, it's not a big deal.  
We can just write all object declarations  
in a single YAML file.  
Well, that's true, but it might make it even  
harder to find the stuff that we are looking for.  
We'd have to continuously search for stuff  
we need to edit in something that could be 25 pages of text,  
at least in multiple files,  
it'd be somewhat organized  
and we'd know we'll find deployment related stuff,  
in the the MySQL deployment of YAML file.  
Enter Helm.  
Helm changes the paradigm.  
Kubernetes doesn't really care about our app as a whole.  
All that it knows is that  
we declared various objects  
and it proceeds to make each of them exist in our cluster.  
It doesn't really know  
that the persistent volume  
and that deployment and that secret  
and that service are all part

of a big application called WordPress.

It looks at all the little pieces that the administrator wanted to have in the cluster and takes care of each one individually.

Helm, however, is built from the ground up to know about such stuff.

That's why it's sometimes called a package manager for Kubernetes.

It looks at those objects as part of a big package as a group, whenever we need to perform an action we don't tell Helm the objects it should touch.

We just tell it what package we want it to act on, like our WordPress app package.

And based on the package name, it then knows what objects it should change and how. And even if there are hundreds of objects that belong to that package.

To make this easier to understand, we'll think about this.

A computer game is contained in hundreds or thousands of files.

There are a few files with the programs executable code, other files with audio, game sounds, and music and other files with graphics, textures, images files with configuration data and so on.

Now, imagine what would happen if we had to download each of these files separately. That would be tedious.

Fortunately, we don't have to go through such horrors as we get a game installer. We run it.

We choose the directory where we want to install, we press install and the installer does the rest, putting thousands of files in their proper location.

Helm does a similar thing and more for the YAML files and Kubernetes objects that make up our application.

Hence, we get advantages like this.

We use a single command to install our whole app, even if it needs a hundred objects.

Helm proceeds to automatically add every necessary object to Kubernetes without bothering us with the details.

We can customize the settings we want for our app or package by specifying desired values at install time.

But instead of having to edit multiple files in multiple YAML files,

we have a single location where we can declare every custom setting.

It's called a value.yaml file,  
and this is where  
we can change the size of our persistent volumes,  
choose the name of our WordPress website,  
the admin password,  
settings for the database engine and so on.  
We can upgrade our application with a single command.  
Helm will know what individual objects need to change  
to make this happen.

We can roll back to the previous  
so-called revision.

We use a single command to uninstall our app.  
It keeps track of all the objects used by each app,  
so it knows what to remove.

We don't need to remember each object  
that belongs to one of our apps anymore  
or use 10 separate commands to remove everything.  
Helm does all the work.

We will look into these commands in more detail  
in the upcoming lectures.

For now, understand that Helm works as a package manager,  
with install or uninstall wizard,  
and also a release manager,  
helping us upgrade or roll back application.

The core thing is that it lets us treat our Kubernetes apps  
as apps instead of just a collection of objects.

This takes a huge burden of our shoulders  
as we don't have to micromanage  
each Kubernetes object anymore.

Helm can do that for us.

∴ Let us now understand the helm concepts.

So back to our WordPress application.

We have now discussed about the challenges.

Let us now see how Helm solves these challenges  
using charts.

Here are the yaml files that we plan to use.

We have the deployment dot yaml  
the secret dot yaml tv dot yaml, PVC dot yaml  
and the service dot yaml.

Now each with its own definition to deploy a component  
of the WordPress application on Kubernetes.

Now we know that some of these have values that might change  
between different environments.

Well, users may prefer to use a different version

of the WordPress image that is used to deploy the WordPress application or different size of disc. And of course, the WordPress admin password is going to be different as well.

So the first step is to convert these files into **templates** where these **values** become variables. The two curly braces indicates that these are variables and that the values specified within our variable names which will be used to fetch these values from another place. So what is that other place where these values are fetched from?

These values are stored in a file named Values dot yaml.

These, this file has the image storage and password encoded variables defined with the values We want these to have.

This way anyone who wants to deploy this application can customize their deployment by simply changing the values from the single file called Values dot yaml.

So a combination of templates plus values dot yaml gives us the final version of definition files that can be used to deploy the application on the Kubernetes cluster. Together the templates and the values file forms a helm chart.

A single helm chart may be used to deploy a simple application like WordPress in our example and it'll have all the necessary template files for different services as well as the values file with the variables.

It also has a chart dot yaml file that has information about the chart itself, such as the name of the chart the charts version, a description of what the chart is some keywords associated with the application and information about the ERs.

Now you can create your own chart for your own application or you can explore existing charts from the **artifact hub@artifacthub.io** and look for charts uploaded by other users.

This hub is called as a repository that stores helm charts.

As of this recording, **there are over 5,700 charts available.**

Search for a chart for the application you're trying to deploy. You

-: Can either search using this web interface or you can search using the command `helm search` followed by `hub` to indicate that you wish to search the Artifact Hub.

The artifact hub is the community driven chart repository but there are other repositories as well such as the Bit nami Helm repository.

To search for charts in other repositories you must first **add a repository to your local helm setup.**

For this run the command `helm repo add` to add the bit NAMI repository with the link to the Bitnami charts repository. Then search the repository using the `helm search repo` command instead of the search hub command. And you can list existing repos using the `helm repo list` command.

Now, once you find the chart the next step is to `install the chart on your cluster`. For this run, the `helm install` command followed by a release name and the chart name.

Now, when this command is run the helm chart package is downloaded from the repository and extracted and installed locally.

So each installation of a chart is called a release and each release has a release name.

So that's the release name that you specify within the `helm install` command.

For example, you can install the same application using the same chart multiple times on a Kubernetes cluster by running the `helm install` command multiple times and each time you run the `helm install` command.

A release is created and each release is completely independent of each other.

Let's proceed to some additional helm commands.

To list installed packages, run the `helm list` command to uninstall packages, run the `helm uninstall` command.

Now, we saw

that we could use the `helm install` command to download and install a helm chart, but if we only need to download it and not install it, then run the `helm pull` command.

Now use the `-u` option because the chart is normally downloaded in a TAR archive format.

The `-u` option will extract its contents after downloading it.

Now once extracted, you can find the contents of the chart.

When you list the files under the extracted folder it's gonna be in the name, same name of the chart and you may open and edit the `values.yaml` file to change any values that if required.

Now, once the changes are made install the local chart using the `helm install` command but by specifying the path to that particular directory.

So that is all that we will be discussing about Helm in this course.

There's so much to learn about Helm so it requires an entire course of its own.

So we are working on an entire course on Helm for beginners

and we'll be releasing that soon.

So watch out for that.

From an exam perspective

we have already covered what is required for you to know.

So head over to the labs and practice working on help.